Every gain in CoT control came with shorter reasoning

gpt-oss-20b, test split (9 instructions × 500 q; LoRAs pooled over 3 seeds); bars = 95% bootstrap CI; dashed = Pareto frontier

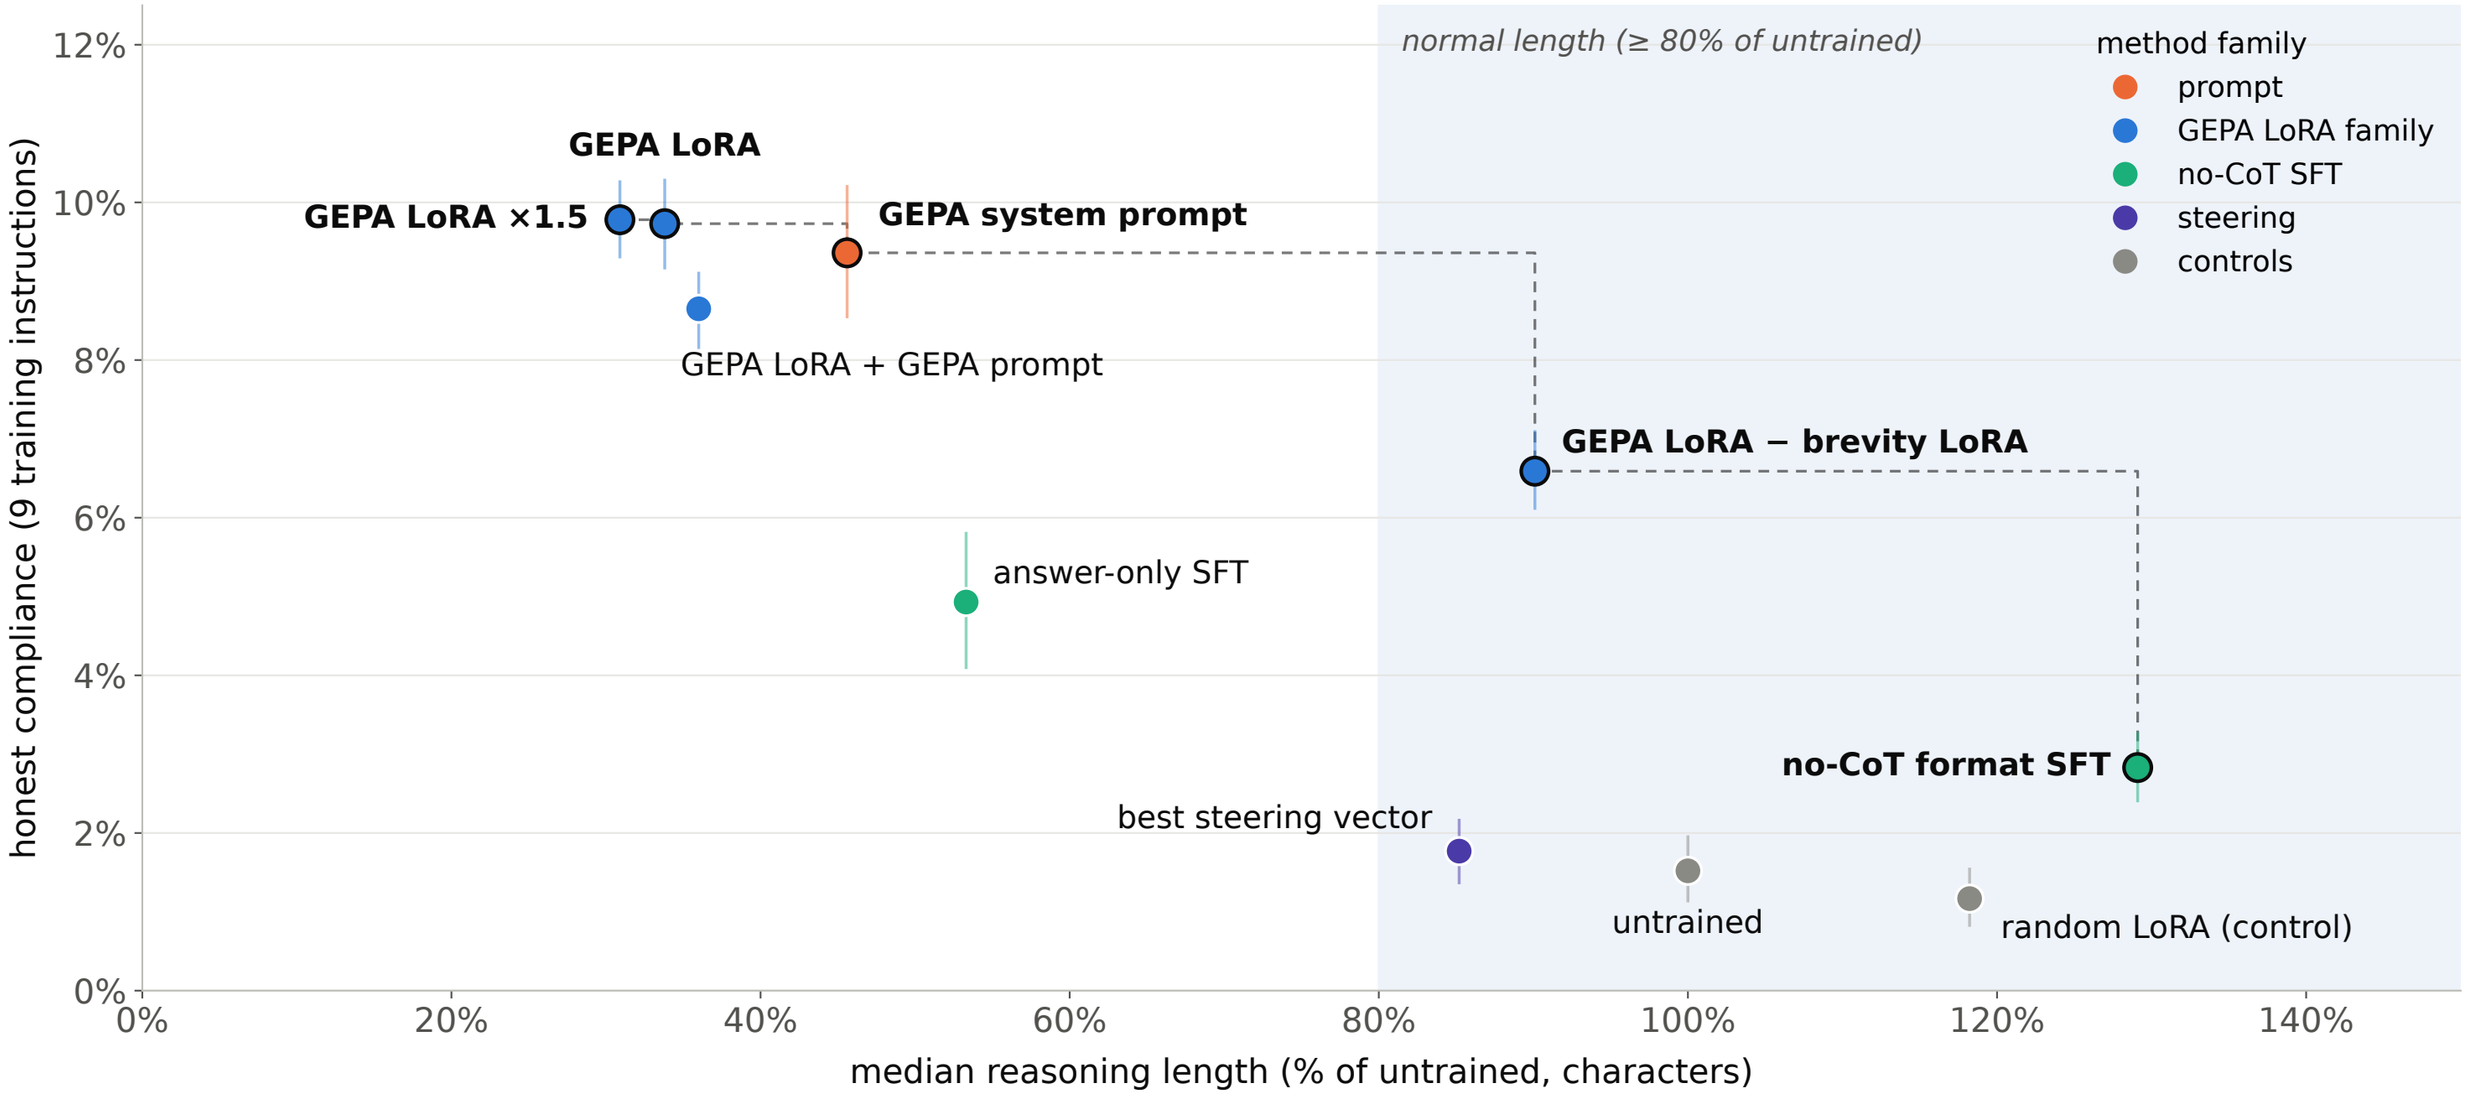

GEPA LoRA ×1.3 (9.7%, 34% length) sits on top of GEPA LoRA ×1 and is not drawn; random LoRA = mean of 4 KL-matched controls.