Control also costs accuracy: every arm above 3% breaks the ±3-point guardrail; the GEPA prompt (−3.6 pts) and the LoRA arithmetic (−4.0) come closest

gpt-oss-20b, test split · honest compliance on the 9 training instructions vs accuracy on 500 unconstrained questions · 95% bootstrap CIs over questions
dashed = Pareto frontier (best compliance at this accuracy or better) · shaded = accuracy within 3 points of the untrained model

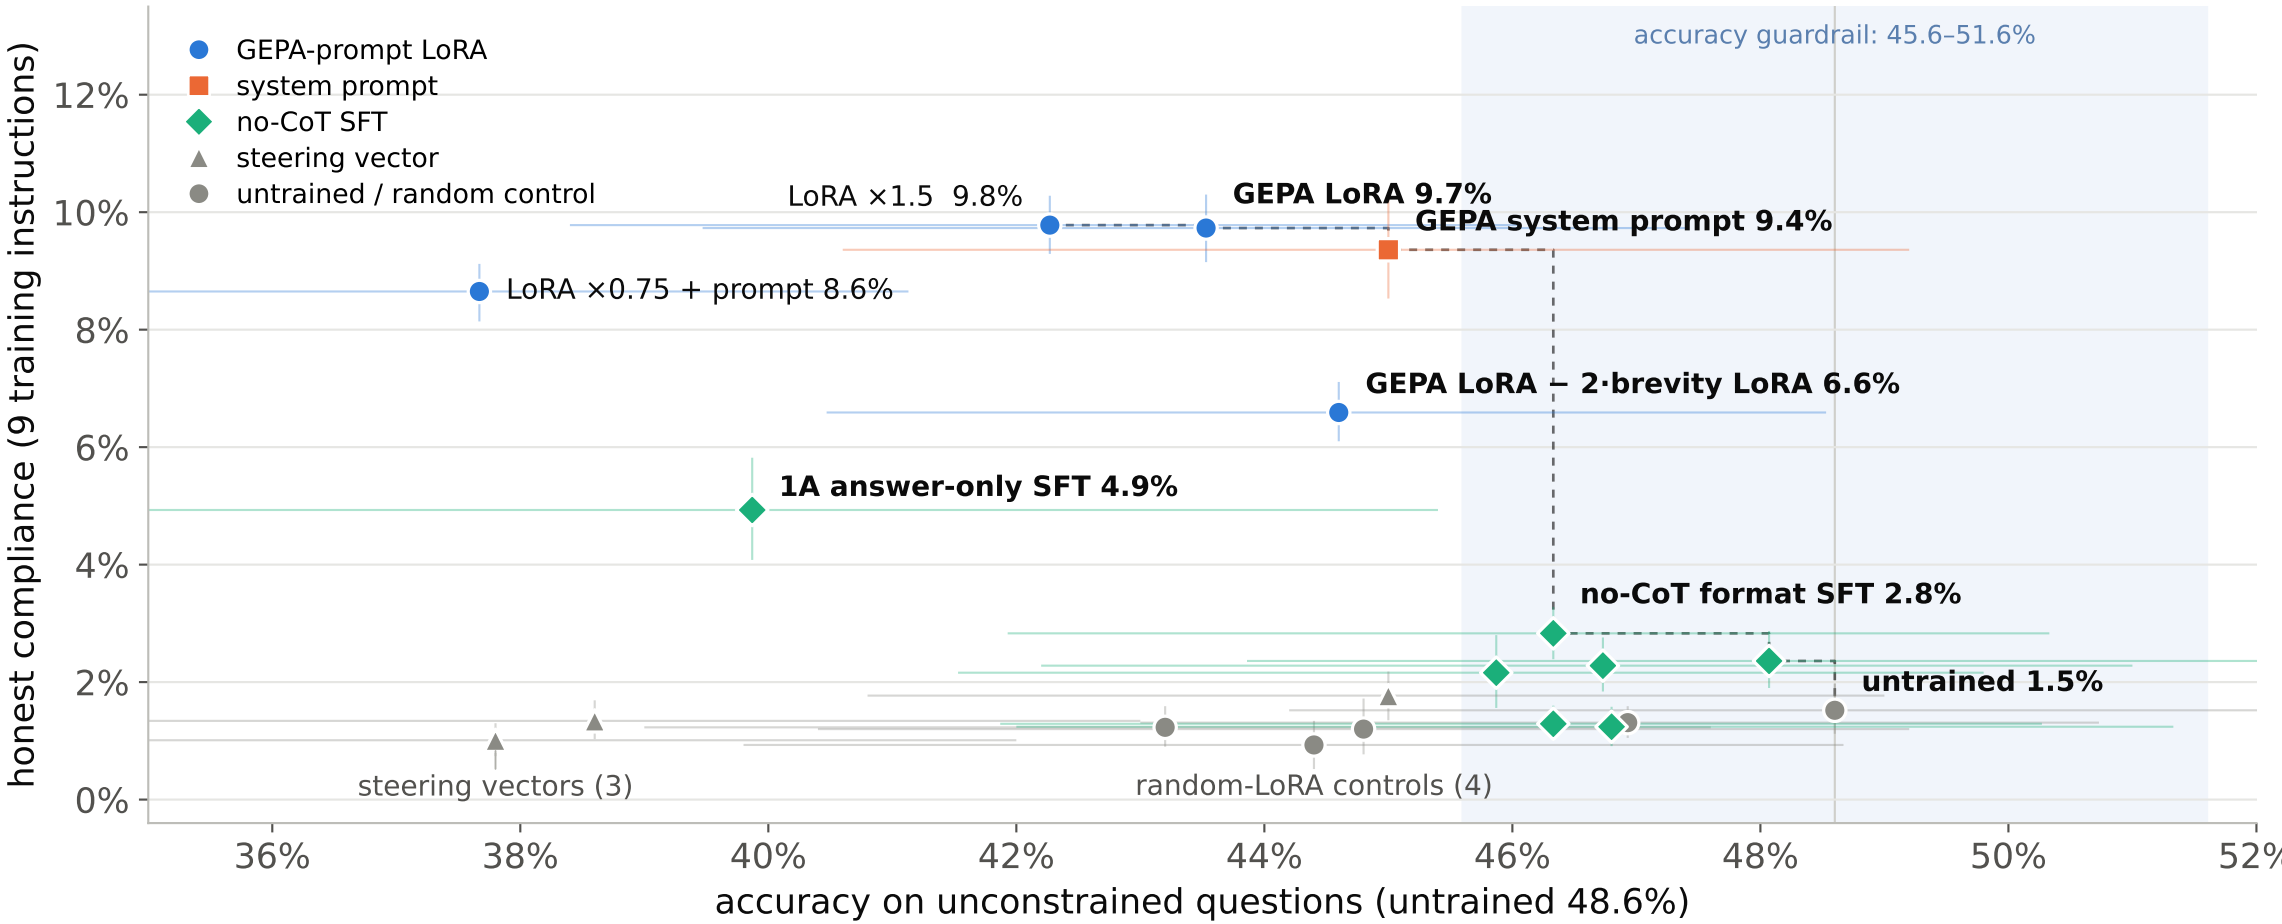

honest compliance = followed the instruction about its reasoning AND answered AND reasoning not empty/degenerate.
Accuracy CIs are ±4–5 points at 500 questions; the pre-registered guardrail compares point estimates.