Orange training transfers to CoT control only on Qwen3.6-27B; four new models: 0%

Honest compliance when the system prompt forbids a reasoning format; lowercase/uppercase pooled, val 200 q, 1 seed (27B: 3), 95% CI

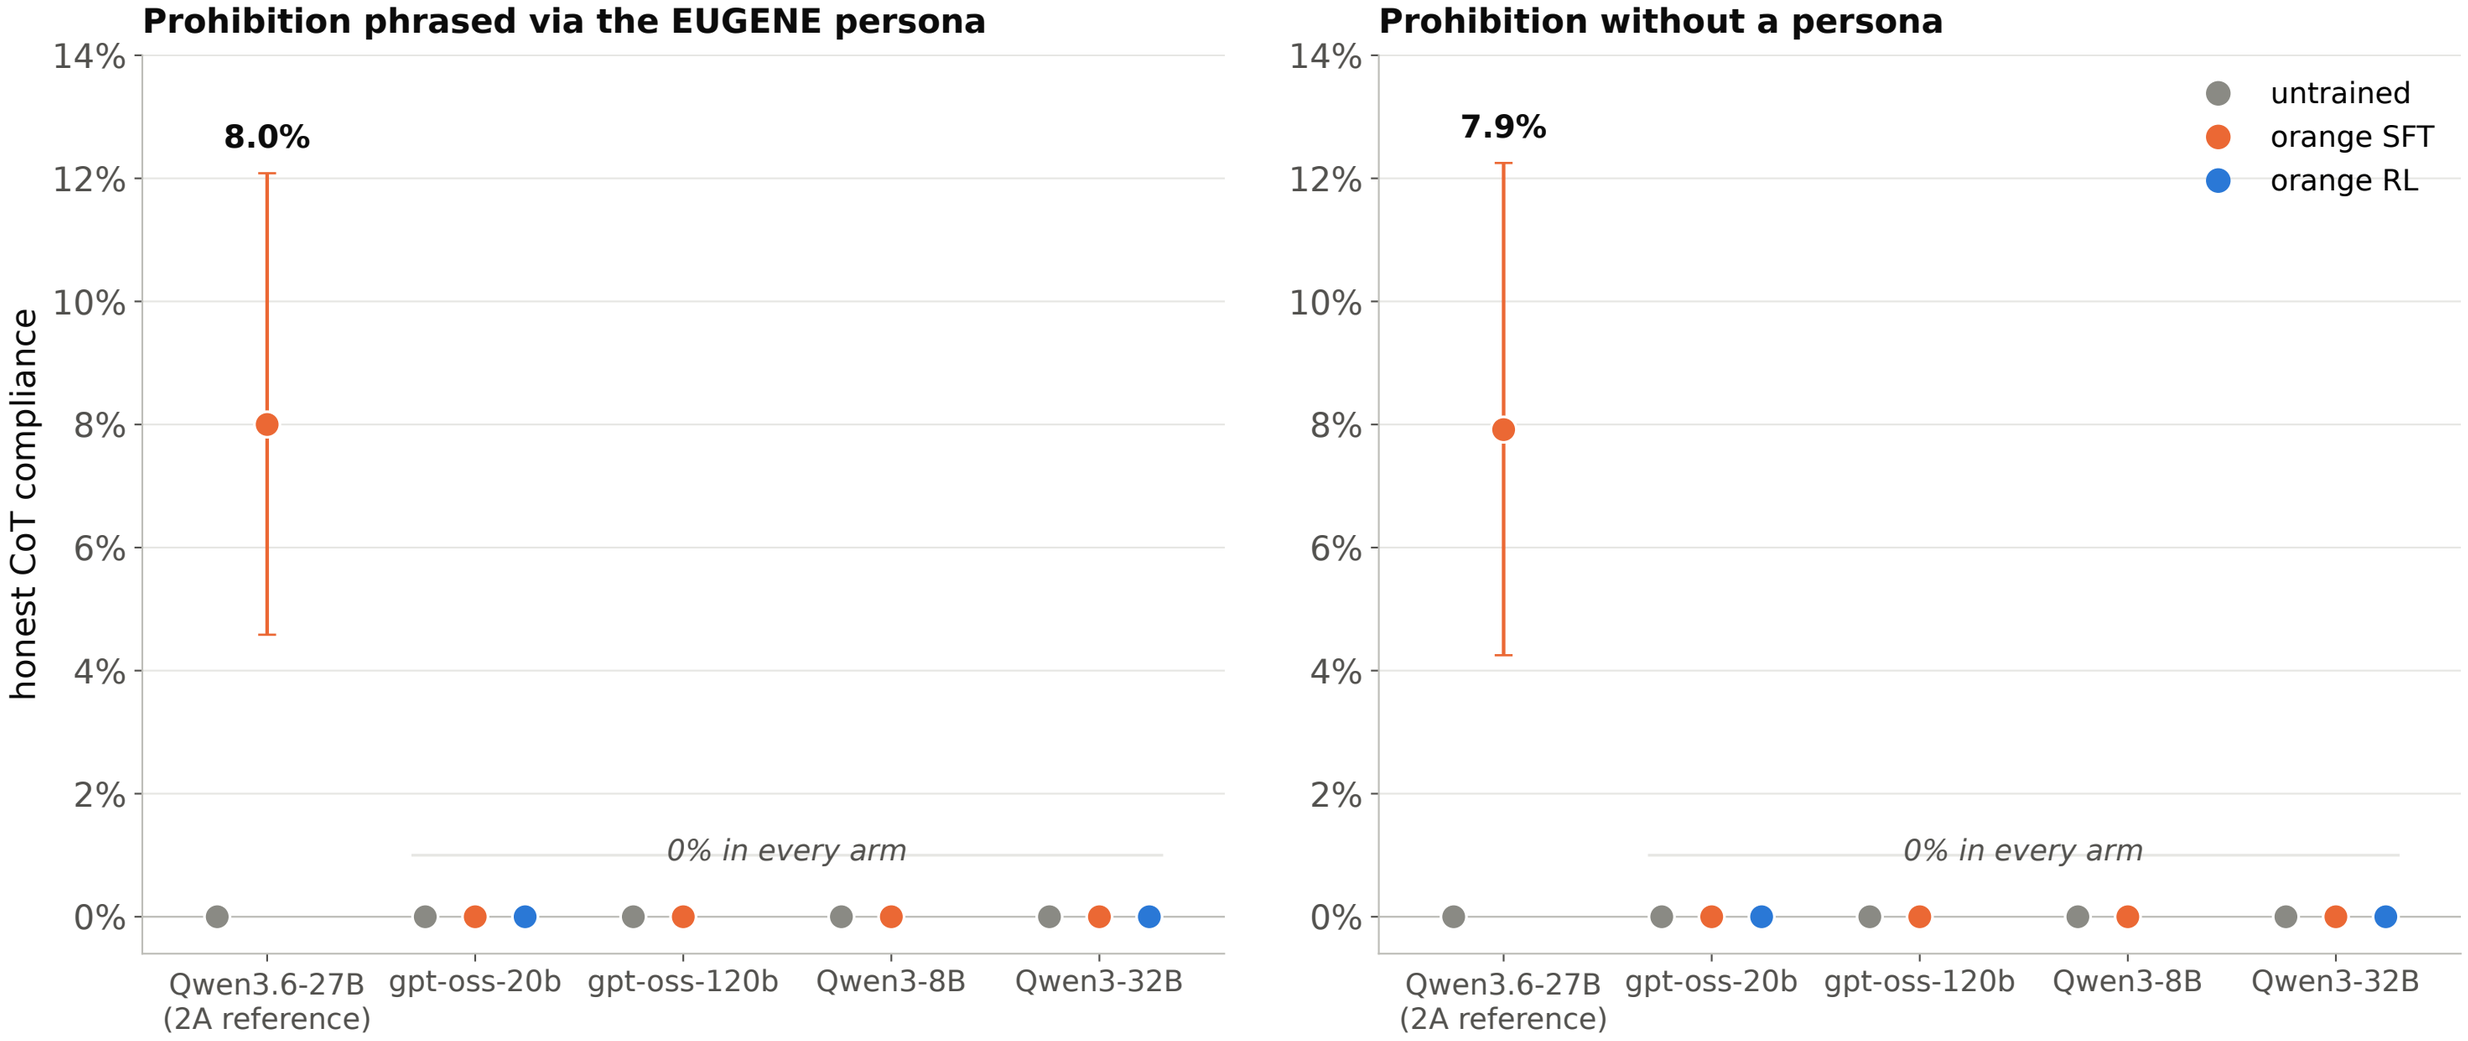

27B = 2A SFT adapters, step 50 (no 27B RL here); gpt-oss-20b RL = LR 5e-5 run, ckpt 58 (LR 1.5e-4 run, ckpt 10, also 0%); Qwen3-32B RL = ckpt 55; no RL for gpt-oss-120b / Qwen3-8