Only generic multiple-choice training without reasoning (1A) beats both the untrained model and a random adapter, and only by 1 to 4 points.

Test split, 'stated' condition, 9 training constraint modes, 500 questions per mode. Whiskers: 95% bootstrap CI over questions and seeds. Text in each panel: 1A's change in honest compliance in percentage points, with its 95% CI.

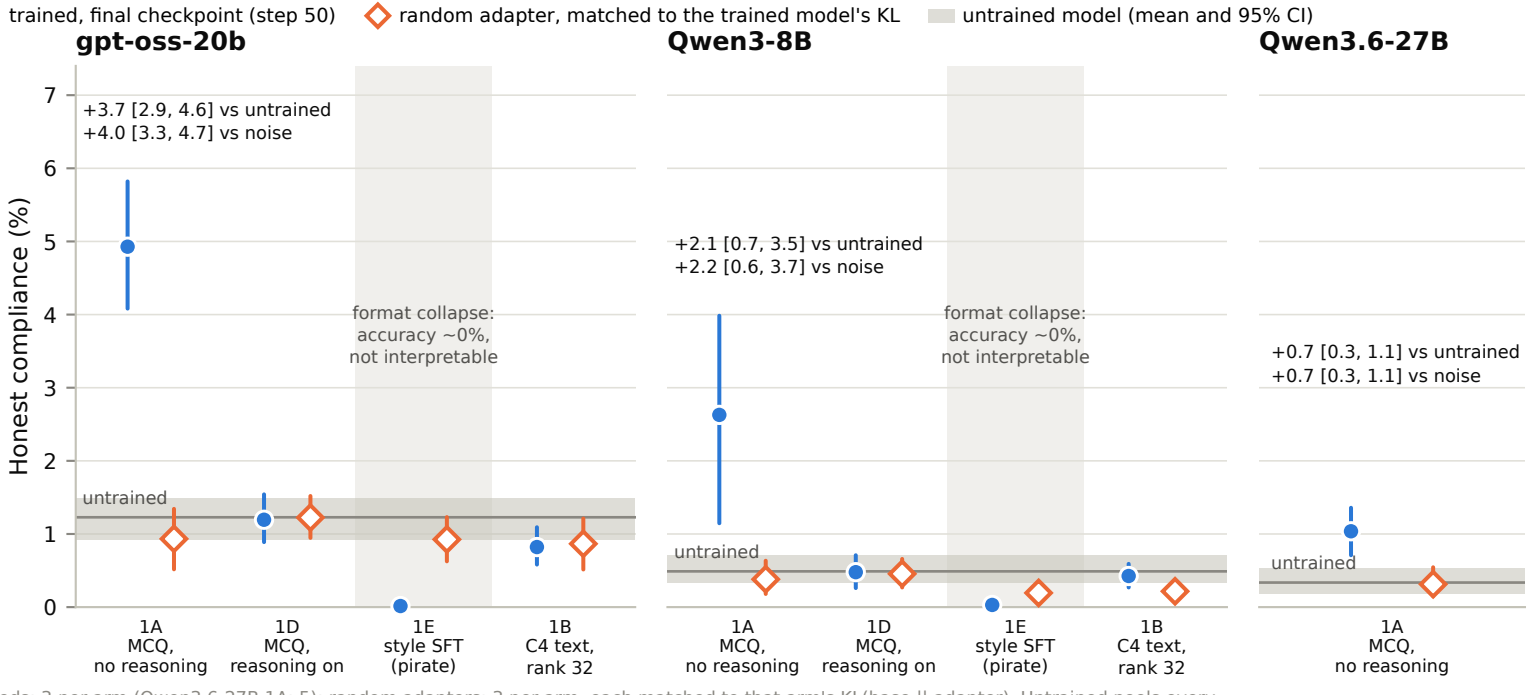

Seeds: 3 per arm (Qwen3.6-27B 1A: 5); random adapters: 3 per arm, each matched to that arm's KL(base || adapter). Untrained pools every arm's checkpoint-0 rollouts for that model (Qwen3.6-27B: 1A only). 1B is at step 50, the same number of steps as 1A. Qwen3.6-27B has no 1B, 1D or 1E runs.