Continued pretraining on web text never raises honest compliance at any adapter size, even when it moves the model as far as 1A does.

Top: honest compliance (test split, 9 training modes, 3 seeds; Qwen3-8B k = 10^3 and 10^4: 2 seeds). Middle: language-model loss on C4 training batches; small adapters fit C4 far less. Bottom: how far the adapter moves the model's output distribution (KL from the base model).

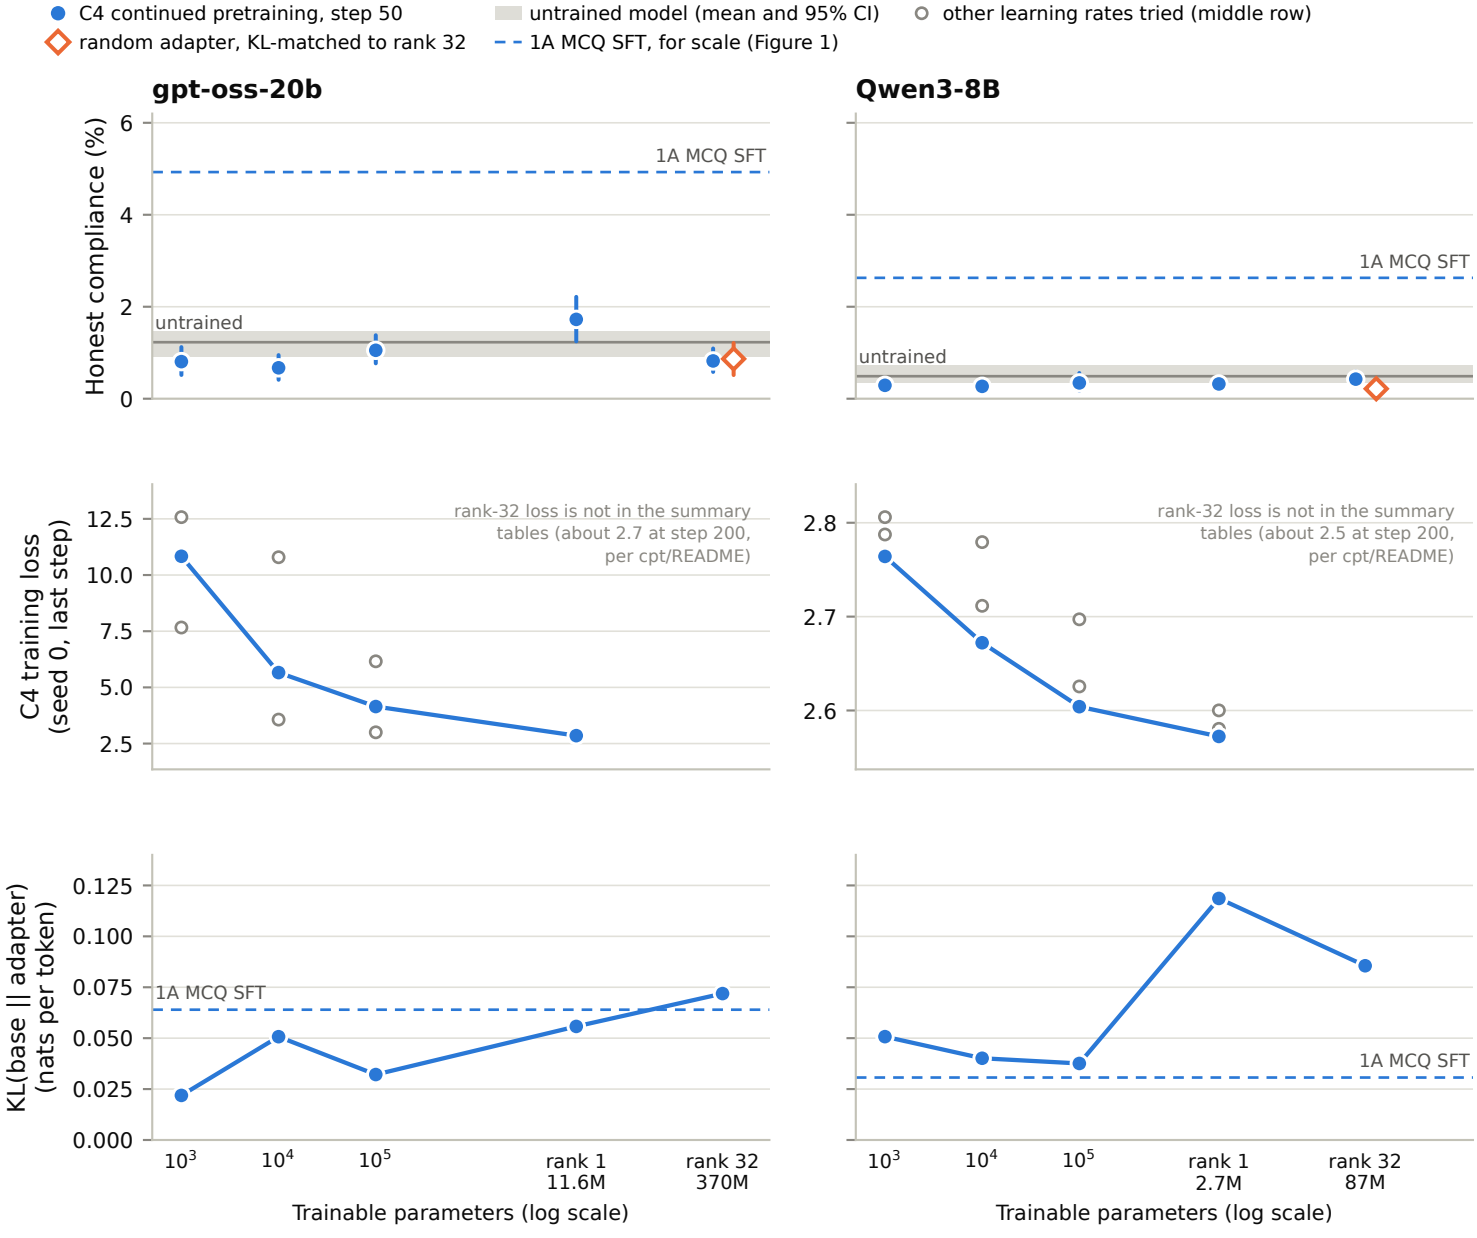

k = number of trainable scalars in a rank-1 LoRA (random subset of its entries); rank 1 / rank 32 = the full adapters. Learning rate per k chosen on val honest compliance with an accuracy guard (cpt/README). Whiskers: 95% bootstrap CI over questions and seeds. Loss: seed 0's final-step training loss. KL: mean over seeds (rank 32: step-50 adapters).