No learning rate lifts honest compliance for any adapter size, including rates high enough to wreck accuracy.

Rank-1 adapters with k trainable scalars, 50 steps of C4, three learning rates per k. Validation split, the four fixed-trait constraint modes. Top: honest compliance. Bottom: accuracy under the constraint, the quantity the learning-rate guard is based on.

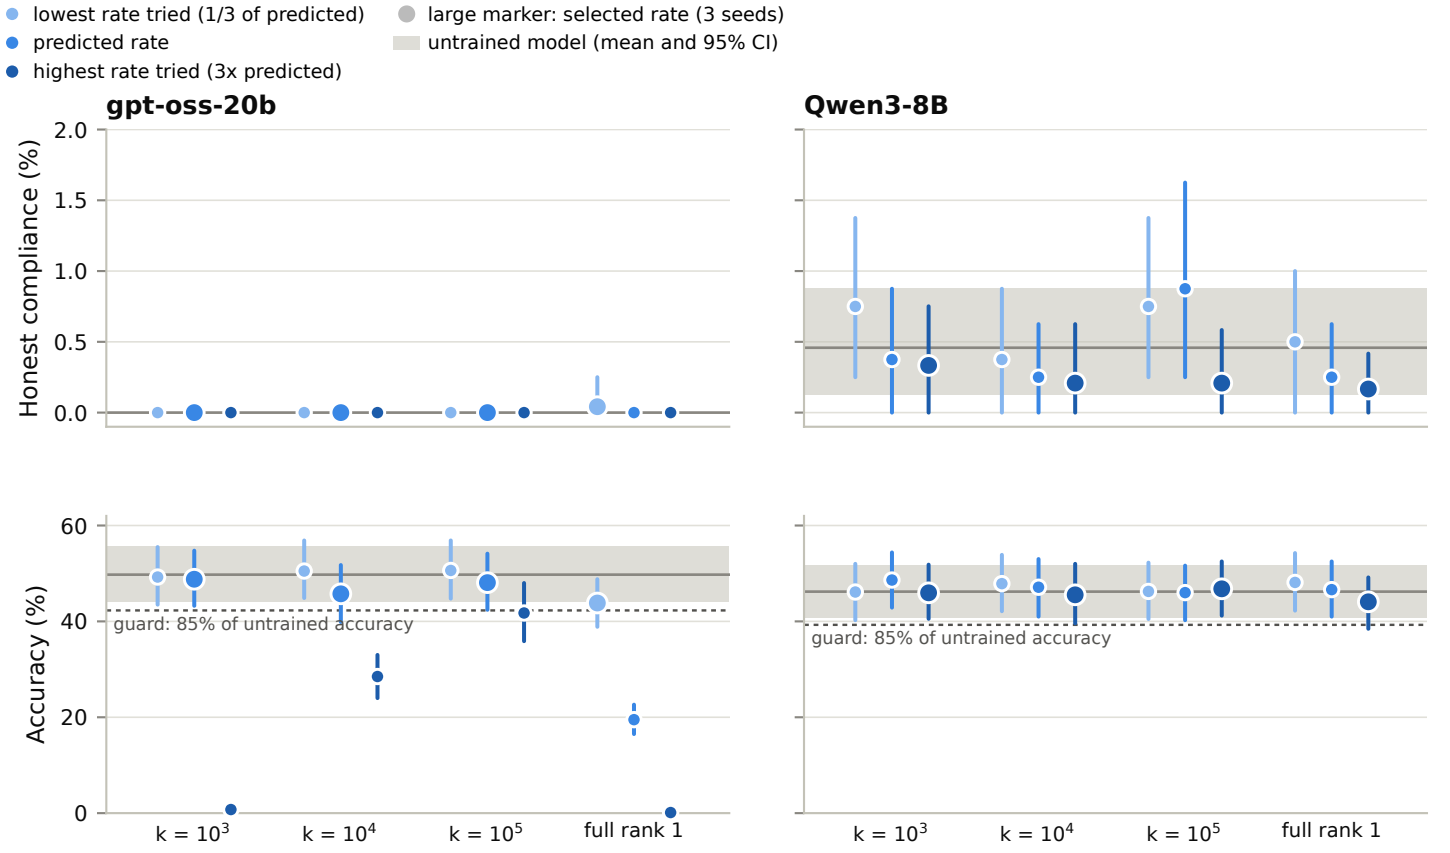

Single seed (800 rollouts) per learning rate, except the selected rate (3 seeds, 2,400 rollouts). Whiskers: 95% bootstrap CI over questions and seeds. gpt-oss-20b is at 0 on these four modes even before training, so this split cannot show a decrease for it; Figure 3 uses the test split.