Full success is near zero for every model, but a graded suppression score separates them: DeepSeek V4 Pro and GLM-5.3 suppress the notation most.

Both panels show the same models and episodes on the same horizontal scale (one percentage point is the same width in each).

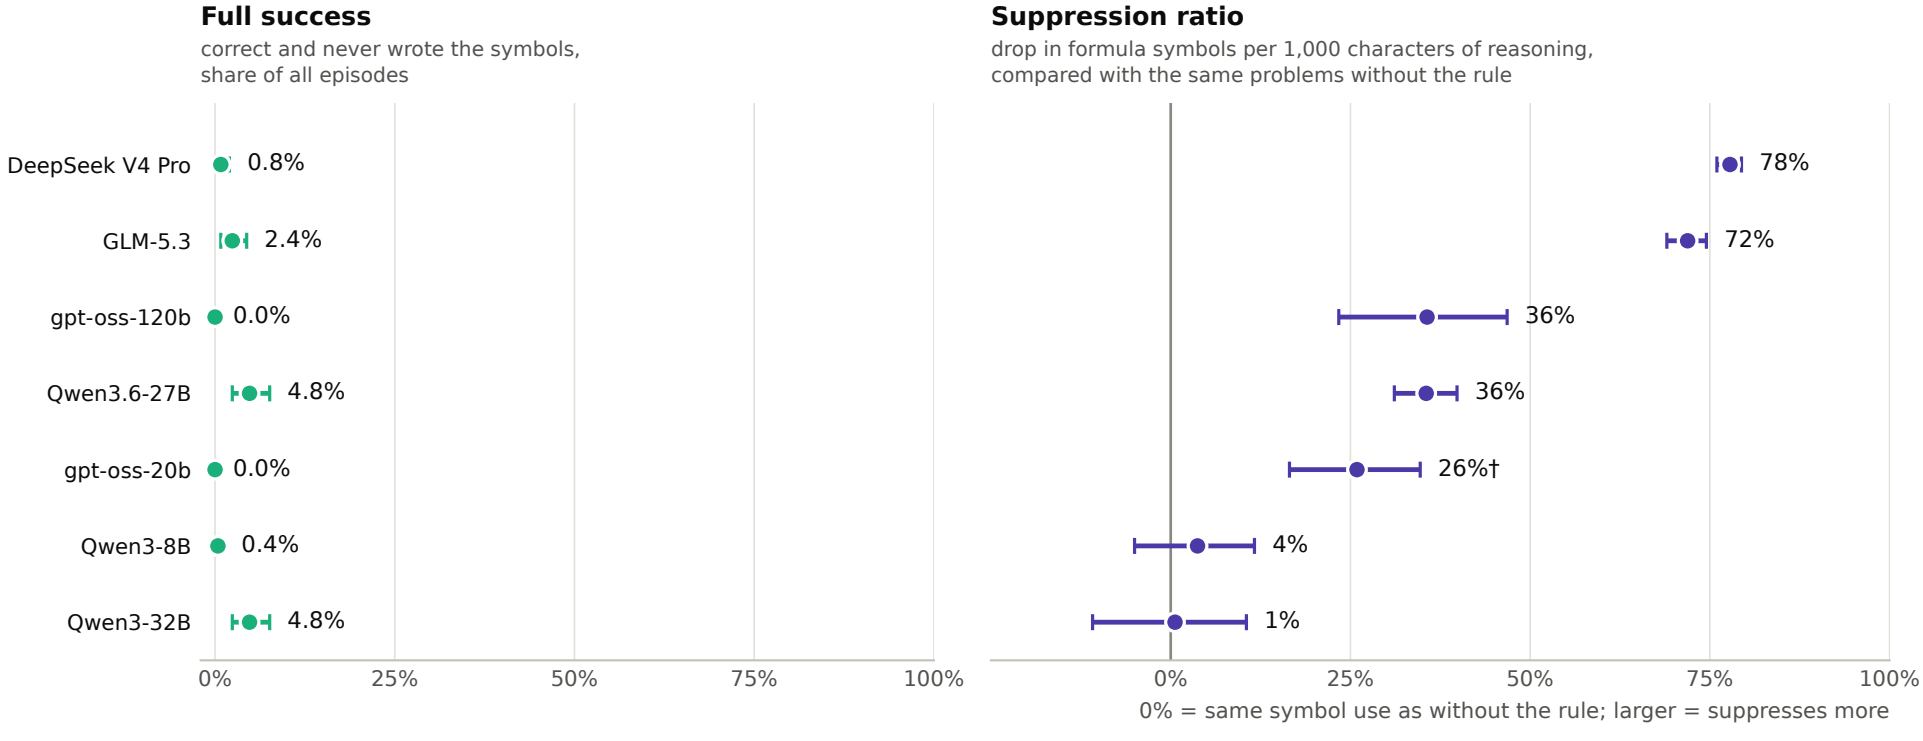

† gpt-oss-20b reads the rule file in only 47% of episodes, so its suppression ratio rests on 51 paired episodes (wide interval). Whiskers are 95% bootstrap intervals over episodes (a 0% rate has a zero-width interval). Arithmetic word problems, 250 episodes per model.