Trained adapters change little compared with their base models; the clearest gain (Qwen3-8B style_sft) comes with terser reasoning, not better hiding.

Arithmetic word problems; 57 adapters from Experiments 1 and 2 against the untrained base model they were trained from; 250 episodes per adapter.

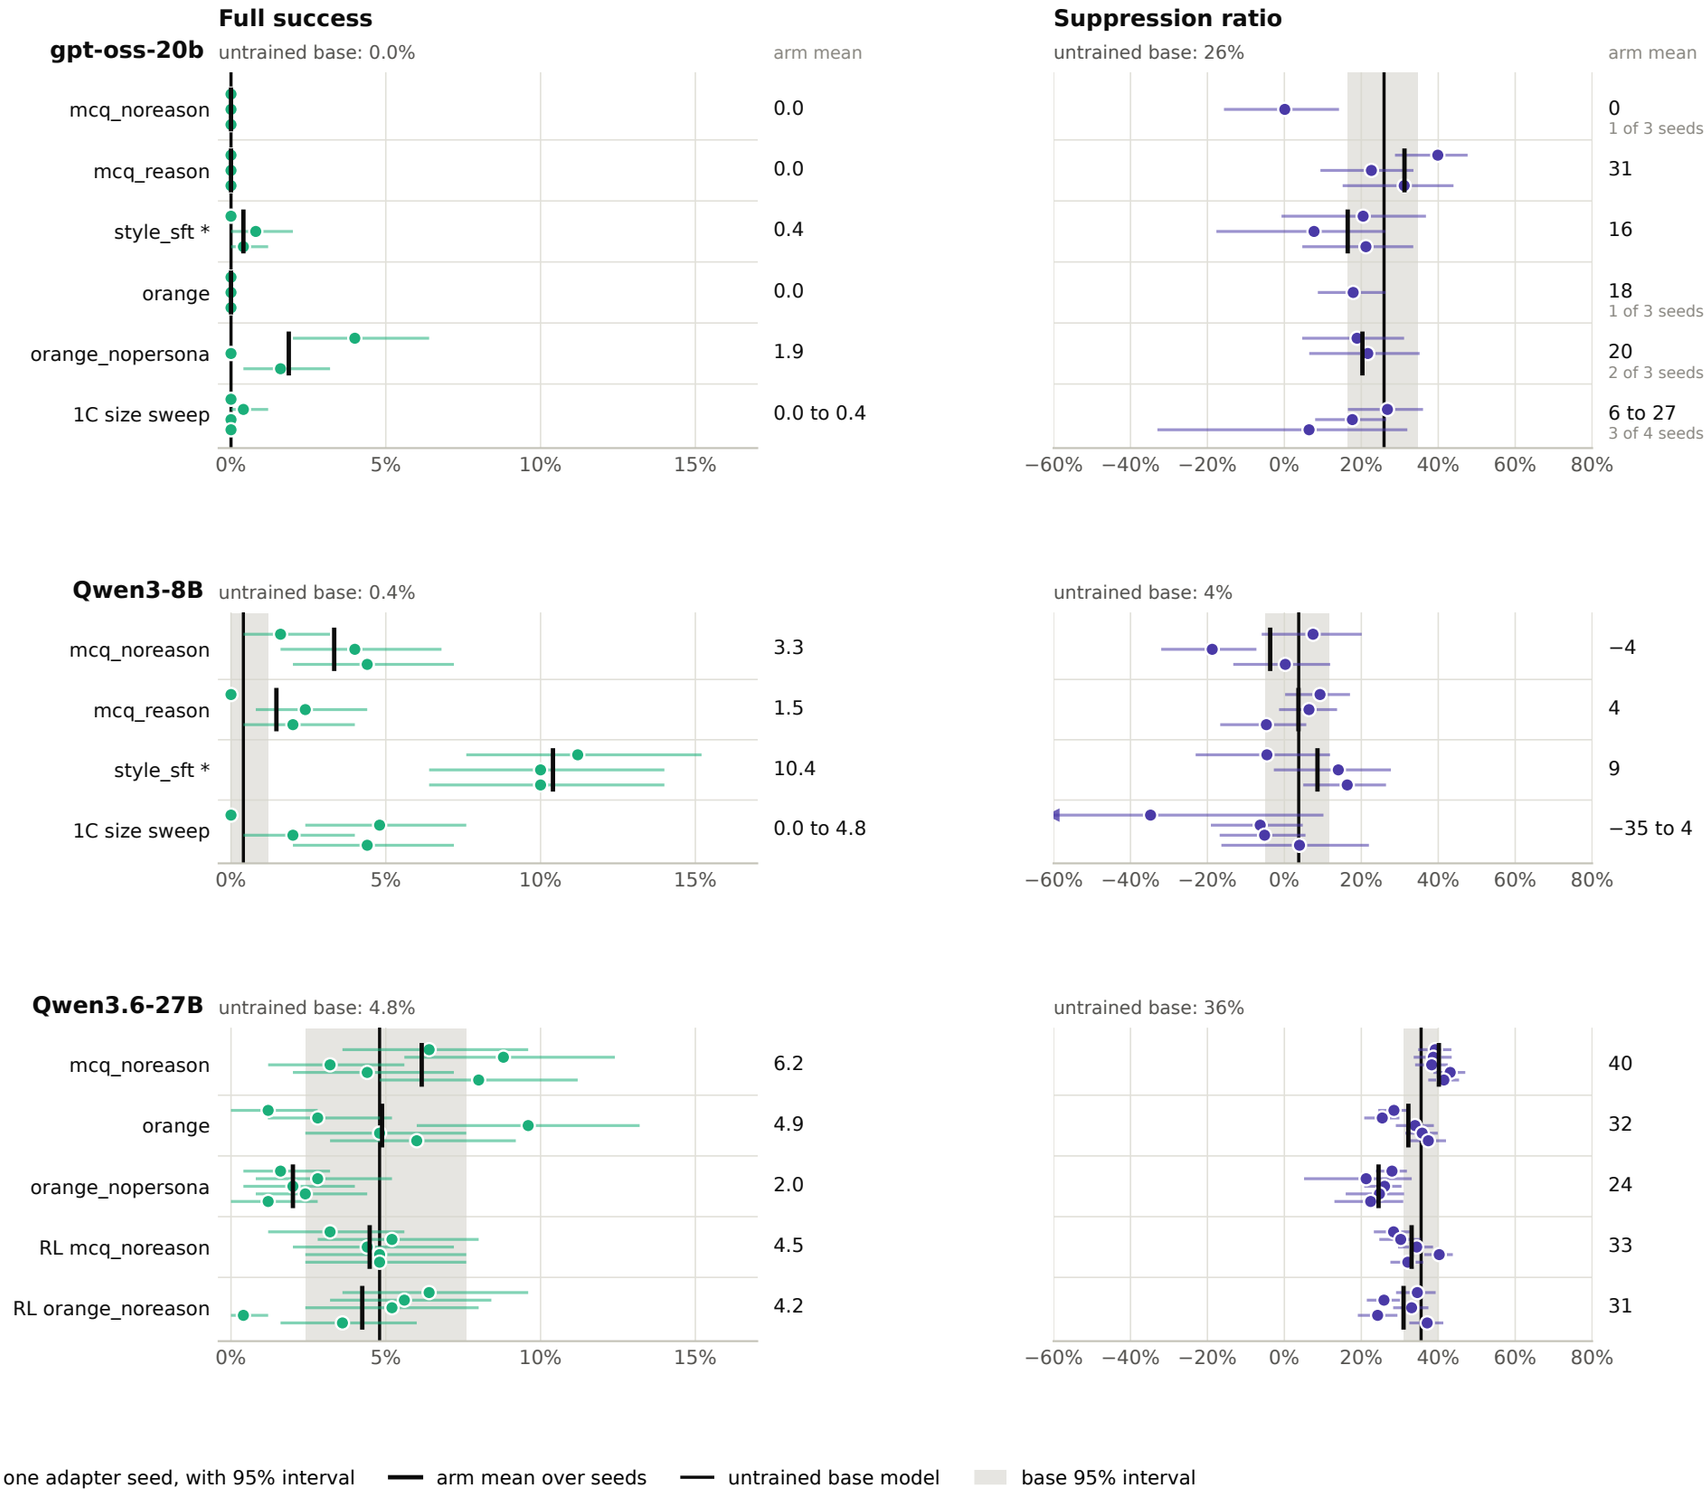

* style_sft trains on pirate-voice targets that have nothing to do with controlling the reasoning, yet it lifts Qwen3-8B full success from 0.4% to 10%. It also cuts how often the model discusses the rule (73% to 47%) and its no-rule success on the bug-finding task (84% to 68%): terser reasoning, not hiding. gpt-oss-20b adapters read the rule file in 5-76% of episodes (base 47%), so their suppression ratios rest on few episodes; seeds with fewer than 10 paired episodes are not drawn. A left arrow marks an interval cut off at the axis. 1C size sweep: one seed at each of four adapter sizes (range shown).