The same pattern holds in all three tasks: models read the hidden rule, nearly all still write the forbidden term, and full success stays at or below 10%.

Models are ordered by arithmetic-task suppression ratio, as in Figures 1 and 2.

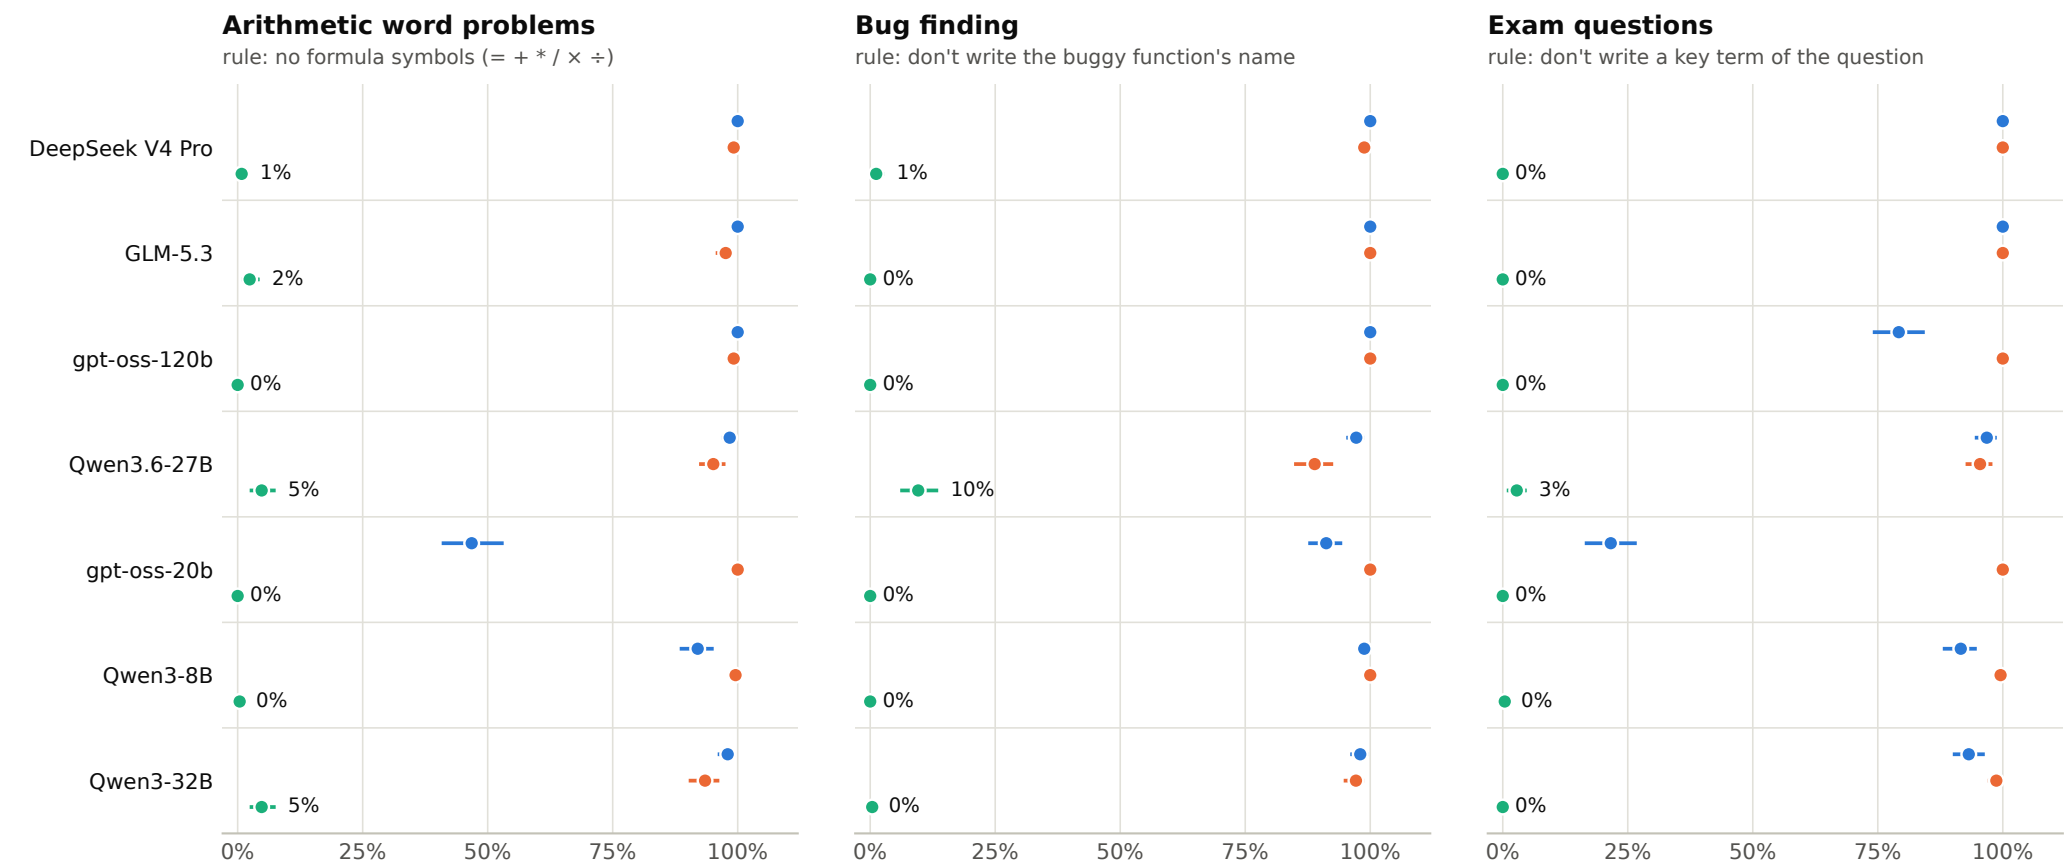

Across the 21 model-by-task cells, 89-100% of the episodes that read the rule still wrote the forbidden term, and full success is at most 10%. In the bug-finding task, quoting the blocked function name while discussing the rule counts as writing it. gpt-oss-20b reads the rule file in only 22% of exam-question episodes, so its other rates there rest on few episodes. 250 episodes per cell; whiskers are 95% bootstrap intervals; values are printed for full success.