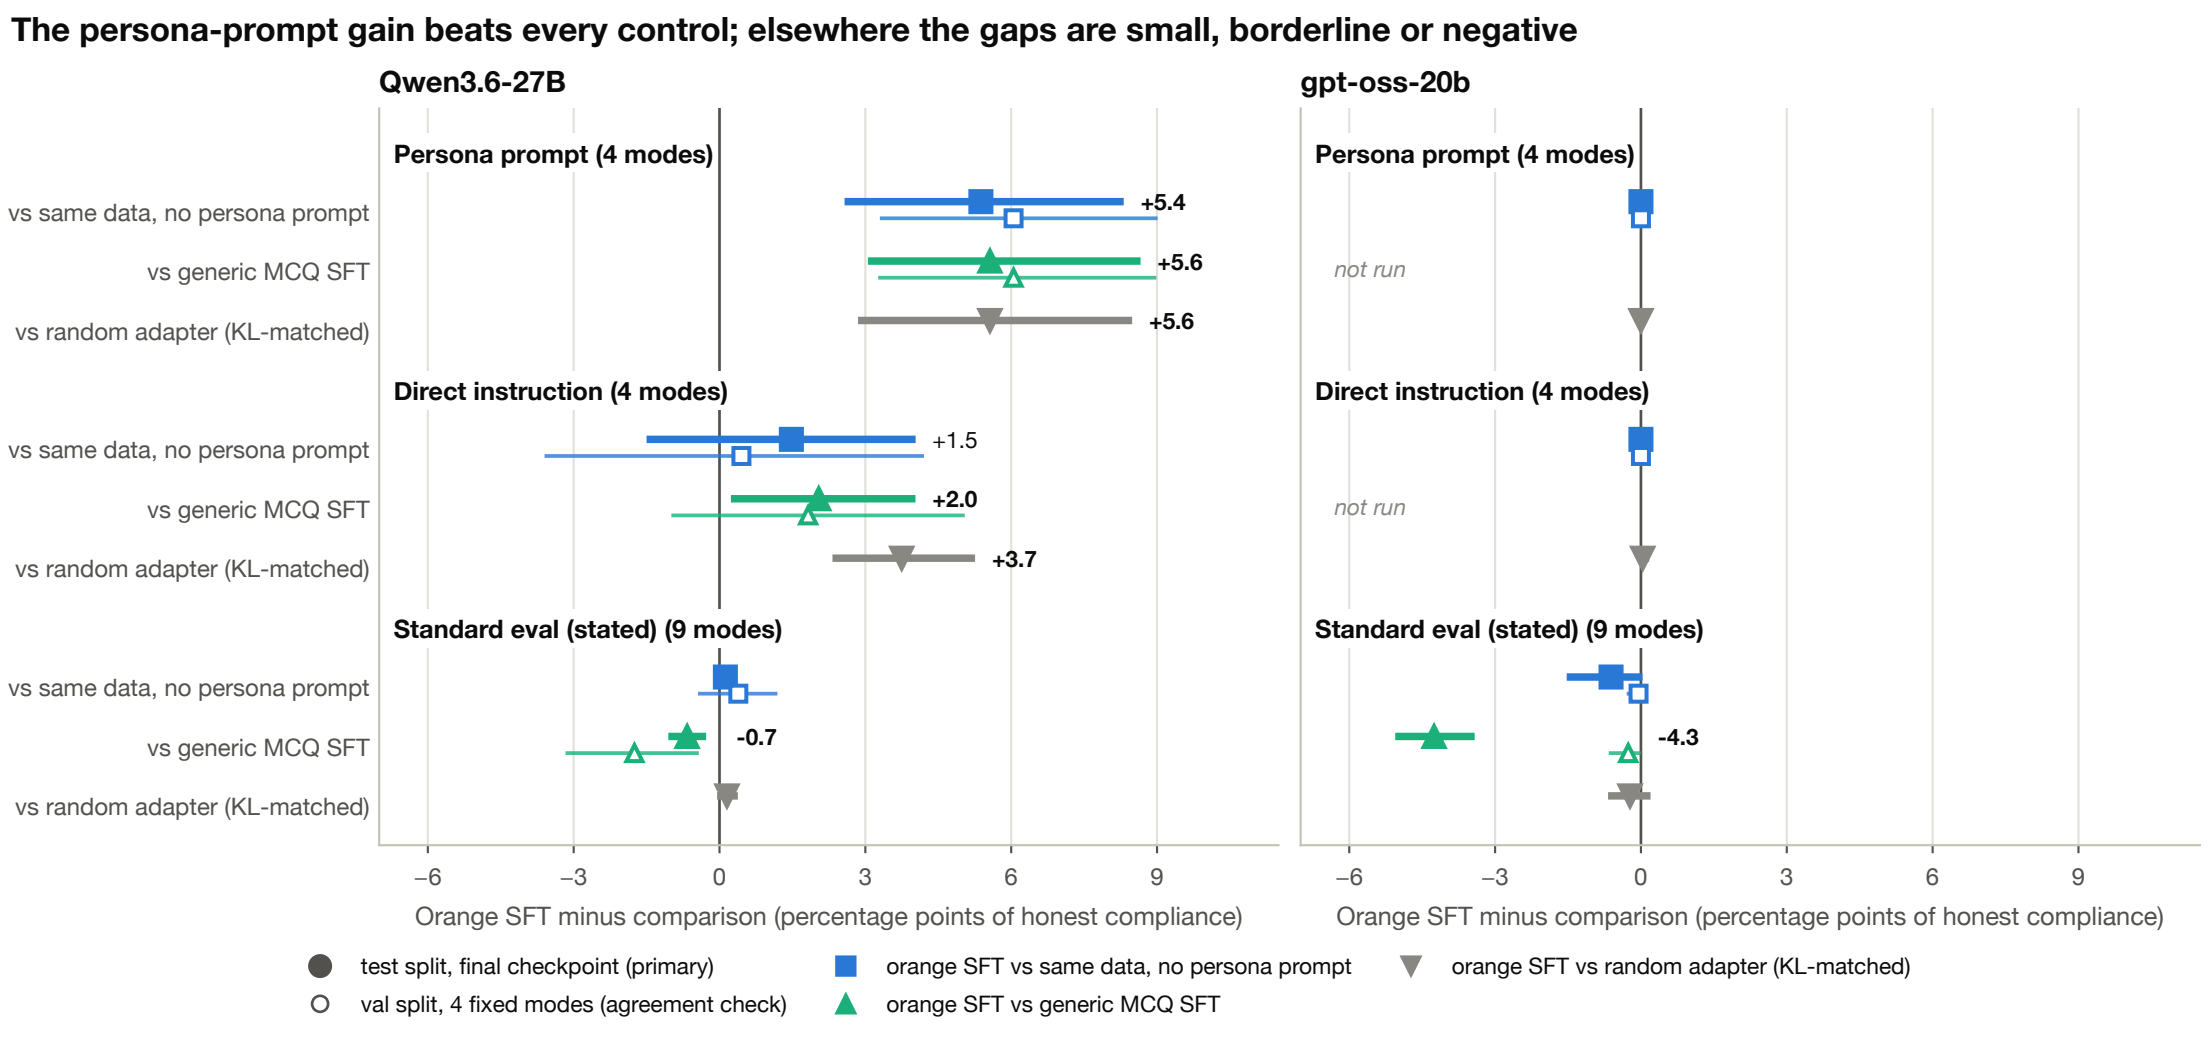
The persona-prompt gain beats every control; elsewhere the gaps are small, borderline or negative
Qwen3.6-27B
gpt-oss-20b
Persona prompt (4 modes)
vs same data, no persona prompt
vs generic MCQ SFT
vs random adapter (KL-matched)
+5.4
+5.6
+5.6
not run
Direct instruction (4 modes)
+1.5
+2.0
+3.7
Standard eval (stated) (9 modes)
-0.7
-4.3
−6
−3
0
3
6
9
Orange SFT minus comparison (percentage points of honest compliance)
test split, final checkpoint (primary)
val split, 4 fixed modes (agreement check)
orange SFT vs same data, no persona prompt
orange SFT vs generic MCQ SFT
orange SFT vs random adapter (KL-matched)