Orange RL leaves honest compliance at the level of the untrained model; only the 2A SFT reference moves it

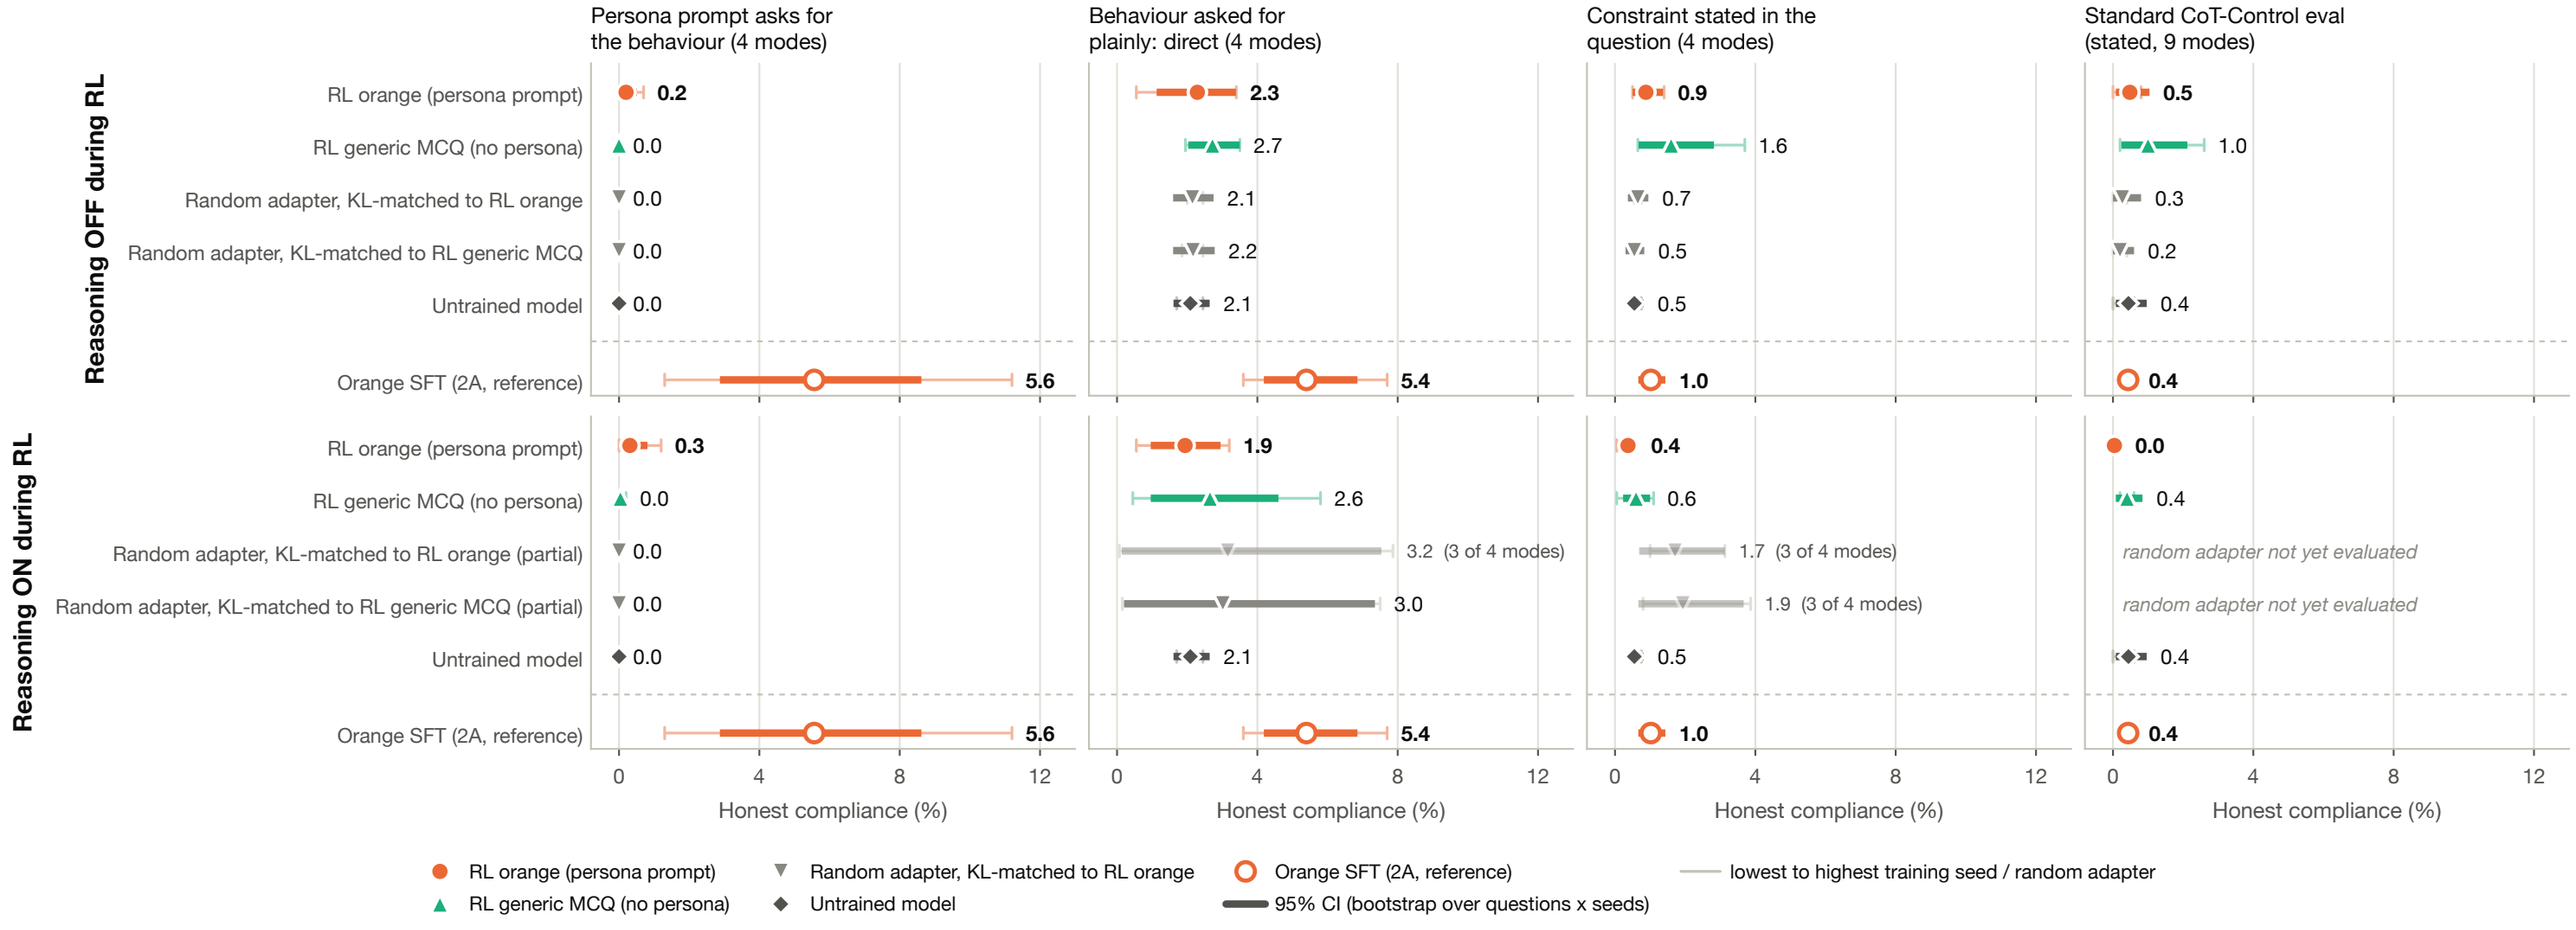

Qwen3.6-27B, test split (500 questions), final checkpoint (step 50), 5 training seeds per RL arm; untrained = checkpoint 0 of the 5 rl_mcq_noreason runs; random adapters: 3 per RL arm, test split.

Honest compliance = followed the stated rule, gave an answer, trace not hollow. 'Standard eval' pools the 9 training modes. Orange SFT (2A) is reasoning-off SFT, 5 seeds, shown for scale only; its standard-eval cell used all 9 modes per question, the RL arms one random mode per question (500 questions per seed).

Partial (faded rows): some random adapters are still being evaluated, so they pool only the modes finished for all 3 adapters (data/fig1_main_result.csv, column `modes`) and are NOT like-for-like with the 4-mode rows next to them (modes differ in difficulty); fig2 panel B compares on identical modes. 'not yet evaluated' = no complete block yet.