Generic-MCQ RL is not smaller than orange RL on the KL axis, but fewer of its prompt groups carried gradient

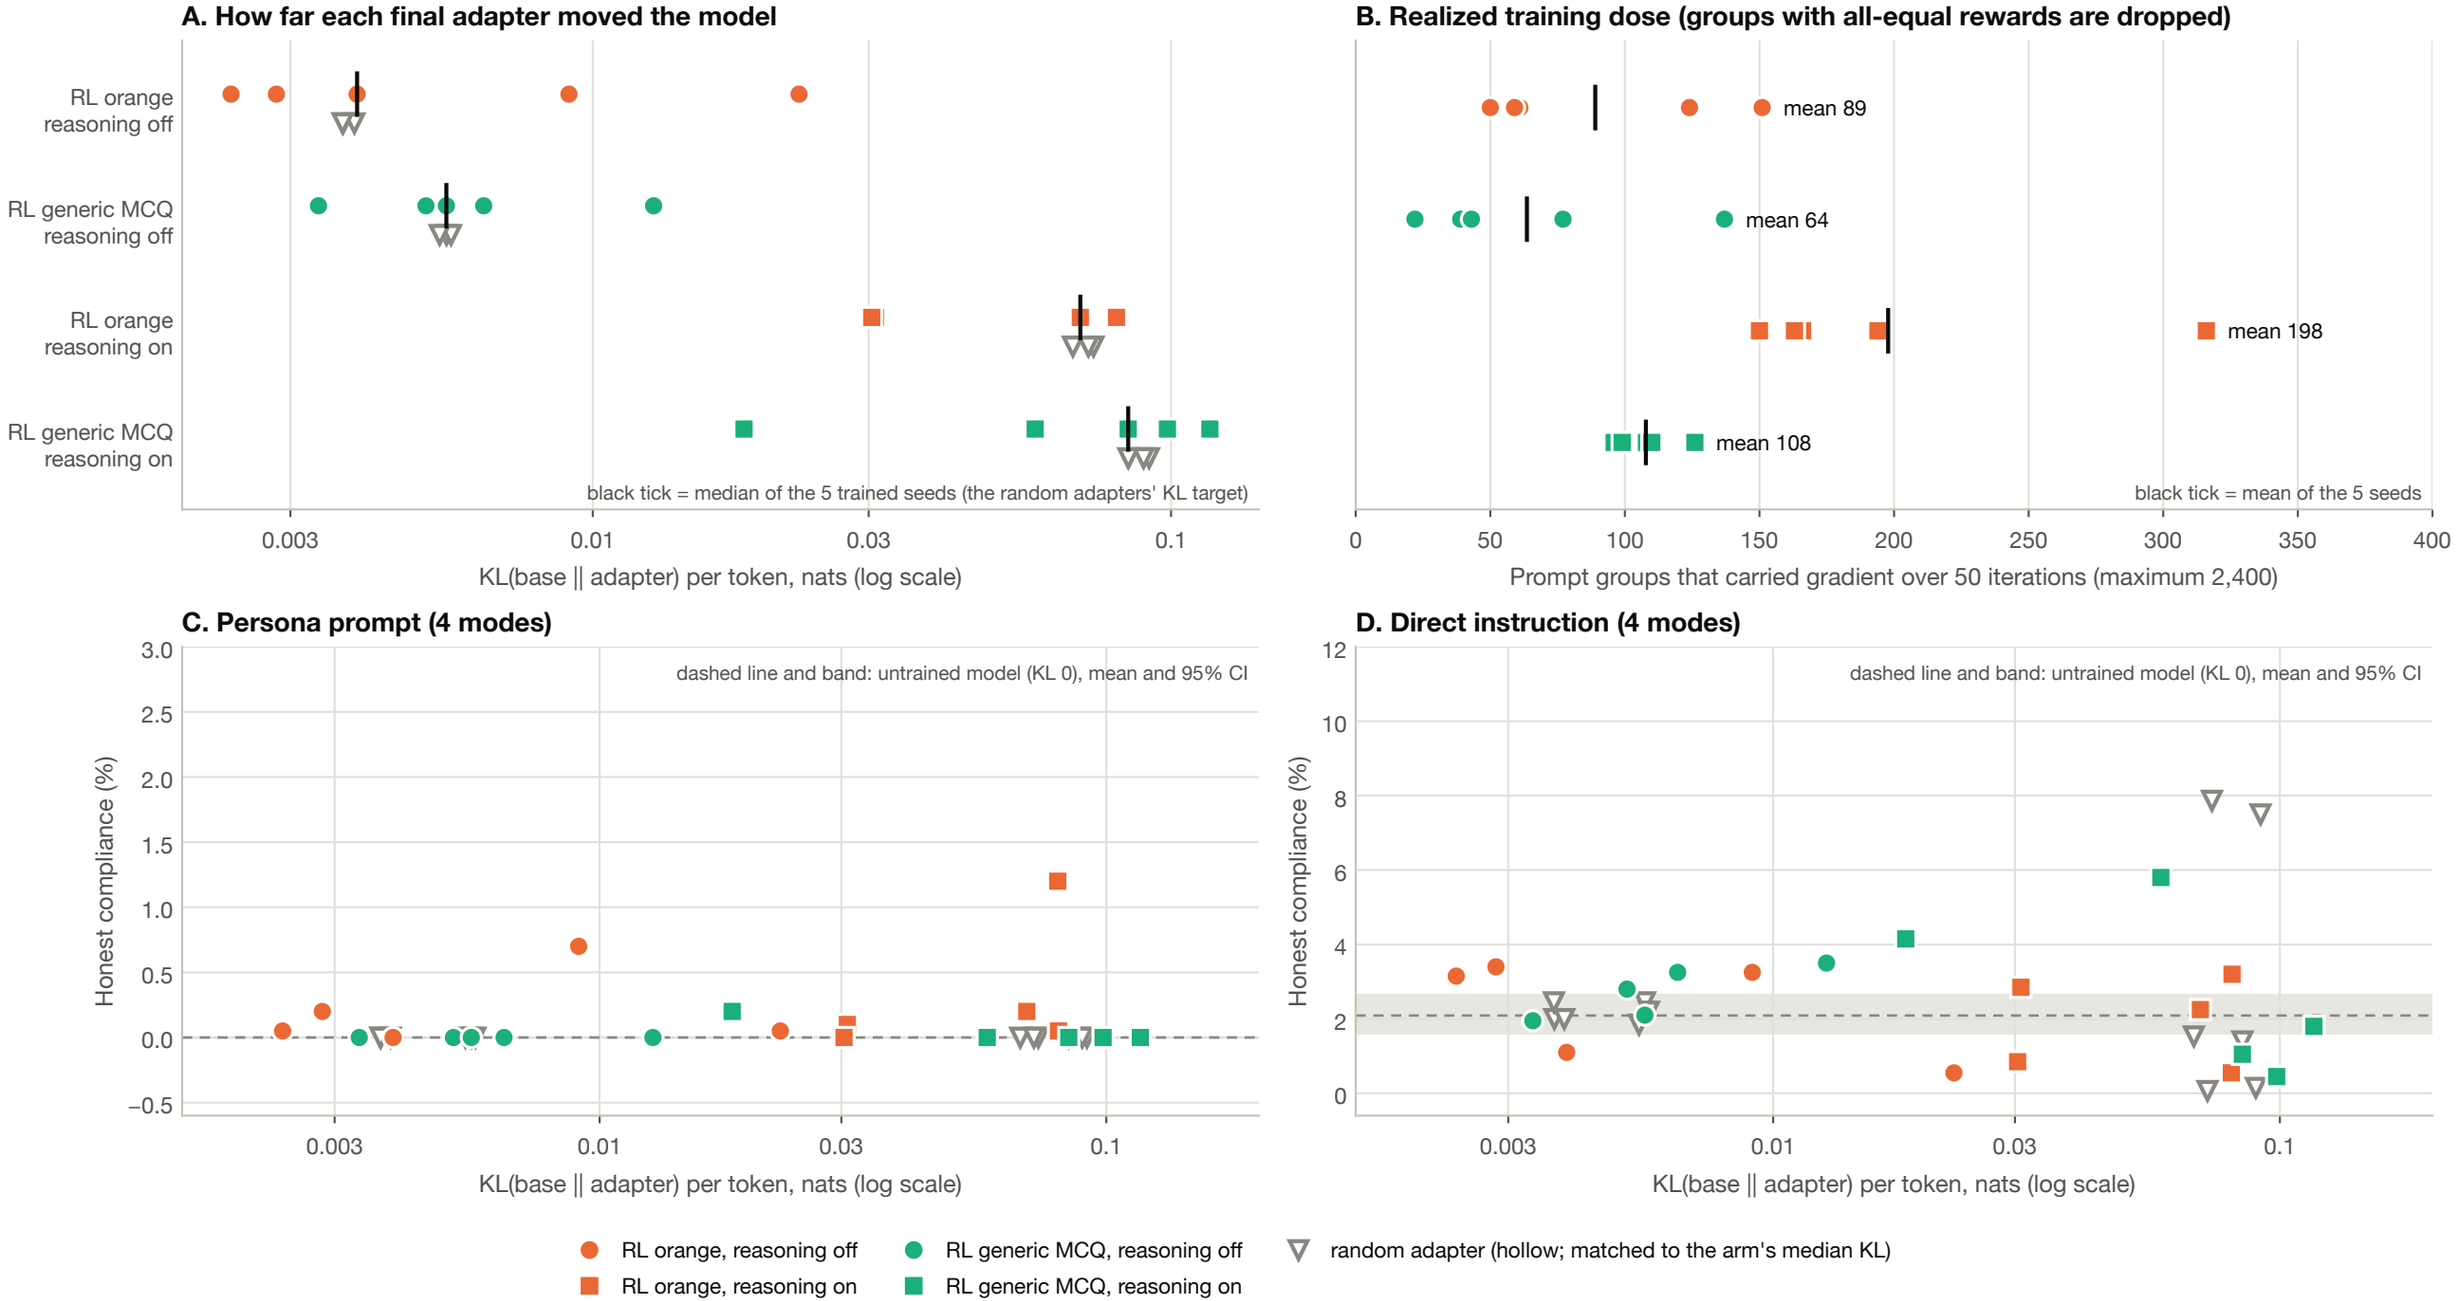

Qwen3.6-27B final adapters (step 50), 5 seeds per RL arm. KL(base || adapter) is the noise agent's measurement on 200 val prompts with T=1 samples from the base model, with no persona prompt, so it may understate how far orange adapters moved the model in their trained context. C and D: one dot per adapter (seed mean), test split, 500 questions x 4 modes; random adapters use only the modes finished for all 3 adapters (fewer than 4 for some reasoning-on arms), so they are not exactly like-for-like with the 4-mode trained dots.