RL shortens reasoning in all 12 arm-by-condition cells; compliance among traces as long as the untrained model's stays at or below 0.1%

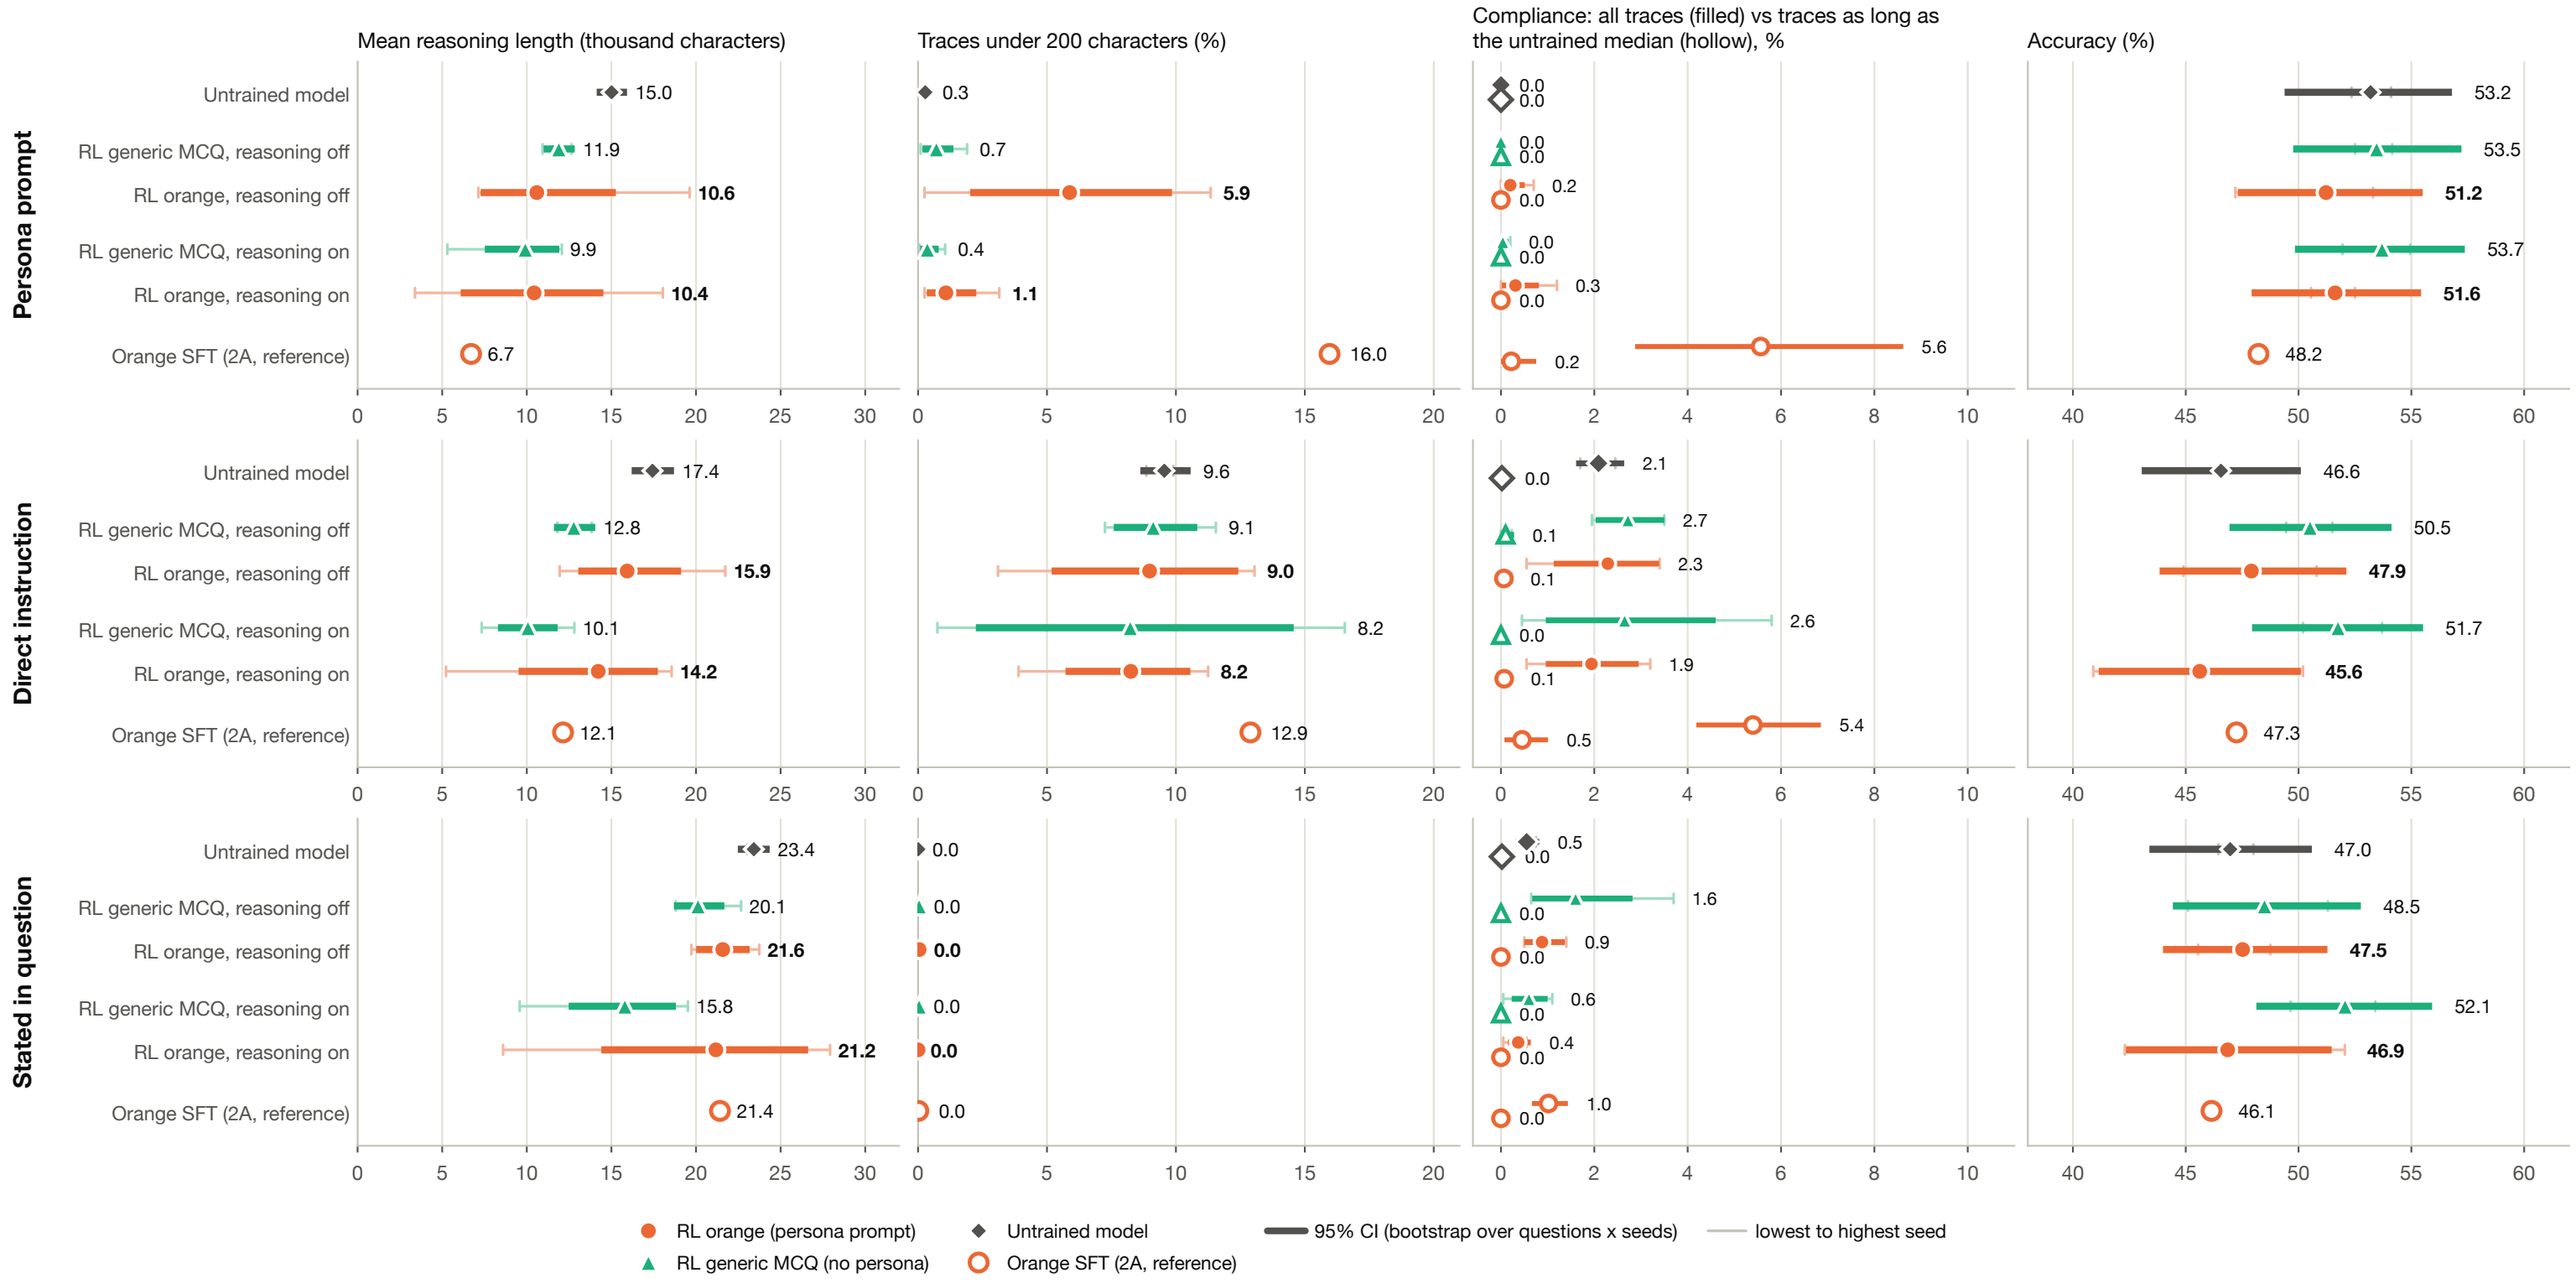

Qwen3.6-27B final checkpoints, test split, 4 fixed modes pooled, 5 seeds per arm, 500 questions. 'As long as the untrained median' = at least the untrained model's median trace length in the same test block; compliance there is strict (raw) compliance, as in 2A. Orange SFT (2A) is reasoning-off SFT, shown for reference; the 2A tables give no interval on length, degenerate share or accuracy.